

Intelligent Agents
Reflection and e-portfolio
Murthy Kanuri
University of Essex
Student ID: 12696139

Main e-portfolio: <https://m-kanuri.github.io/> (m-kanuri, 2024, updated 2025)

Module e-portfolio: <https://m-kanuri.github.io/Module6.html>

Table of Contents

| | |
|--|------------------|
| <i>Introduction.....</i> | <i>3</i> |
| <i>WHAT – Description of Experience.....</i> | <i>3</i> |
| <i>SO WHAT – Analysis and Interpretation</i> | <i>4</i> |
| <i>NOW WHAT – Future Actions and Application</i> | <i>5</i> |
| <i>Conclusion</i> | <i>6</i> |
| <i>References</i> | <i>6</i> |
| <i>Appendix A: e-portfolio artefacts and repositories</i> | <i>7</i> |
| <i>Appendix B: Unit 6 Team E Project Report</i> | <i>10</i> |
| <i>JIRA Board</i> | <i>10</i> |
| <i>GitHub Repository</i> | <i>11</i> |
| <i>Appendix C: Selected e-portfolio evidence</i> | <i>12</i> |
| <i>Collaborative Discussions</i> | <i>12</i> |
| <i>Module Wiki</i> | <i>13</i> |
| <i>Creating Parse Trees</i> | <i>14</i> |
| <i>Deep Learning in Action</i> | <i>15</i> |
| <i>GitHub Repository for Individual Project</i> | <i>16</i> |
| <i>SQLite Screenshot</i> | <i>17</i> |
| <i>Appendix D: Personal Skills Matrix and Action Plan (PDP)</i> | <i>18</i> |
| <i>Appendix E: Skills Matrix.....</i> | <i>19</i> |

Introduction

This reflection documents my learning experience in the Intelligent Agents module. I worked with my team to build the Intelligent Forensics Agent, designed to automate file identification, process the data, and securely archive specified file types for analysis. In this reflection, I will critically analyse my collaborative development experience, the technical knowledge I acquired, and how the module influenced my understanding of intelligent systems. Using Rolfe et al.'s (2001) reflective framework, I outline *what* occurred during the module, *so what* these experiences meant for my professional and personal growth, and *now what* actions I will take to apply this learning in future contexts.

WHAT – Description of Experience

The module covered core intelligent-agent concepts: reactive, deliberative, and hybrid architectures, agent communication, Natural Language Processing (NLP), and adaptive methods, including Artificial Neural Networks (ANNs).

In Unit 6, I coordinated the team by setting up our communication channels, structuring tasks in JIRA, and chairing regular check-ins. I supported the digital forensics analysis, reviewed and organised the GitHub repository, and contributed to the final report and peer reviews. This sharpened my planning discipline and showed how clear roles, lightweight processes, and visible timelines lift overall quality.

We faced practical challenges in coordinating across time zones, merging our work on GitHub, and agreeing on consistent performatives (ask-if, tell, perform). Content-based identification (e.g., comparing python-magic with file-extension heuristics), safe cross-platform scanning (psutil), and a read-only, auditable flow became key design anchors. Learning about Hybrid Architectures (Unit 4) and Communication Languages (Unit 5) clarified how deliberative reasoning can complement reactive behaviour even in a primarily single-agent system so that quick responses do not block longer-running goals.

Throughout the module, the collaborative discussions (agent-based systems, agent communication languages, and Deep learning in action) allowed me to analyse topics such as hybrid architectures and the advantages and limitations of agent communication languages like KQML

(Knowledge Query and Manipulation Language). Later units (e.g., parse-tree ideas in Unit 8 and deep-learning concepts in Unit 10) showed how systems parse inputs and adapt from data, which I related back to our pipeline's identify-and-archive stages.

SO WHAT – Analysis and Interpretation

Working on the digital-forensics agent had a tangible emotional and intellectual impact. Early on, I felt anxious about the project's breadth—making file-type detection reliable, keeping scans read-only and safe across platforms, and integrating these pieces into a coherent discover-identify-decide-act-audit pipeline. Iterative sprints, small spikes, and peer feedback reduced this anxiety. As our tests passed and the SQLite-based metadata trail and modular parsers held up, I moved from uncertainty to curiosity to ownership. One key lesson is that we waited too long to build automated test pipelines.

I found my learning strongly echoed Kolb's (1984) experiential cycle. With every sprint and code review, I moved through the phases of concrete experience, reflection, conceptualisation, and experimentation, making the process much more active. The discussions and GitHub pull requests significantly boosted my confidence. I learned how to communicate technical ideas succinctly, negotiate design decisions effectively, and accept feedback without defensiveness, all vital behaviours for professional teamwork.

The Module Wiki in Unit 4 reinforced a shared understanding. Contributing to and reading succinct, referenced entries on key terms (e.g., hybrid agents, performatives, provenance) encouraged independent research and collective meaning-making. These concepts informed my decision to keep a reactive message handler for responsiveness while a deliberative module processed goals in the background.

Academically, the module deepened my grasp of agent-based computing, especially its focus on autonomy, social ability, and proactivity (Wooldridge, 2009). Our project bridged theory and practice: speech-act theory (Searle, 1969) helped me reason about message intent; KQML highlighted the value of explicit semantics but also the engineering trade-off versus simple method calls; and readings on semantic alignment

(Payne & Tamma, 2014) reminded me that interoperability and shared meaning are as critical as functional correctness. We ultimately favoured lightweight, explicit message schemas over full ACL (Agent Communication Language) complexity sufficient for our scope without over-engineering.

Our team dynamics closely followed Tuckman's (1965) stages. The 'Storming' phase was initially stressful, with disagreements over coding priorities and code structure. Adding a shared branching model, a Pull Request (PR) template, and a brief "definition of done" (tests + documentation + lint) cut down on merge conflicts in the next sprint. As the team grew, I noticed that trust and communication standards improved, which greatly benefited both the speed of our project and the quality of our code.

Ethically, studying AI incidents and governance (Nasim et al., 2022) sharpened my view that even seemingly mundane automation (like file discovery) needs accountability: provenance and audit trails, least-privilege access, clear scope boundaries, and transparent reporting. These considerations influenced our read-only default, metadata logging, and preference for content-based over name-based identification.

Finally, in Unit 11, I extended the work individually by putting the Unit 6 design in practice. By focusing on end-to-end execution and adding unit tests, I was able to solidify my learning. This process quickly revealed gaps such as lower test coverage or issues with edge-case handling that I promptly addressed through small refactors and clearer documentation.

NOW WHAT – Future Actions and Application

This module has reshaped how I learn and deliver. Reflection is no longer a post-hoc report but a continuous improvement habit. In future projects, I will schedule short reflective checkpoints at sprint boundaries and tie them to concrete actions (e.g., a test to add, a doc to clarify, a risk to retire).

Technically, I plan to explore agent-based simulation (e.g., Mesa) to test coordination strategies and to prototype reinforcement learning for adaptive behaviours where safe and justified. Exposure to NLP and deep learning prompts me to investigate how compact language interfaces (not

full ACLs) can make agent interactions more natural while preserving auditability.

Professionally, I will continue using GitHub to document contributions transparently and practise ethical engineering principles (provenance, minimum necessary data, clear user consent). Our teamwork underlined the value of clear documentation, equitable task allocation, and empathy in distributed settings habits I will carry into real-world AI projects.

Ultimately, the module reinforced that effective intelligent systems combine technical rigour with ethical reflection, and that continual learning is essential in a fast-moving field.

Conclusion

The Intelligent Agents module moved me from theoretical interest to practical capability. I learned that teamwork, critical thinking, and ethics are as important as code quality and that small, well-evidenced design choices (read-only defaults, metadata trails, modularity) compound into trustworthy systems. I attended all seminar sessions led by the module tutor (Samuel Danso), which strengthened my academic grounding and prepared me to contribute responsibly to intelligent systems in practice.

The appendix includes my e-portfolio details and covers e-portfolio elements and formative activities.

References

- Finin, T., Weber, J., Wiederhold, G., Genesereth, M., Fritzson, R., McKay, D., McGuire, J., Pelavin, R., Shapiro, S. & Beck, C., 1994. KQML as an agent communication language. *Proceedings of CIKM*, pp. 456–463.
- **Kanuri, M., Espag, J., Kirwan, F. and Jittipattanakulchai, G. (2025) *Intelligent Forensics Agent – Group E Design Report (Unit 6)*. University of Essex Online. (= 2025a)**
- Kanuri, L.M., 2025. *Intelligent Agents – Unit 6 Development Team Project Report*. University of Essex Online. Available at: <https://m-kanuri.github.io/artefacts/IA-Unit06->

[DevelopmentTeamProjectReport.pdf](#) [Accessed 17 October 2025]. (= 2025b)

- Kanuri, L.M., 2025. *Intelligent Agents – Unit 11 Development Individual Project Presentation*. University of Essex Online. Available at: <https://m-kanuri.github.io/artefacts/IA-Unit11-Development%20Individual%20Project-%20Presentation.pdf> [Accessed 17 October 2025]. (= 2025c)
- Kolb, D.A., 1984. *Experiential learning: Experience as the source of learning and development*. Englewood Cliffs, NJ: Prentice-Hall.
- m-kanuri (2024, updated 2025) *Personal e-portfolio*. GitHub Pages. Available at: <https://m-kanuri.github.io/> (Accessed: 17 October 2025).
- Nasim, S.F., Ali, M.R. & Kulsoom, U., 2022. Artificial Intelligence Incidents & Ethics: A Narrative Review. *International Journal of Technology, Innovation and Management (IJTIM)*, 2(2), pp. 52–64.
- Payne, T.R. & Tamma, V., 2014. Negotiating over ontological correspondences with asymmetric and incomplete knowledge. *Autonomous Agents and Multi-Agent Systems*, 13(1), pp. 517–524.
- Rolfe, G., Freshwater, D. & Jasper, M., 2001. *Critical reflection in nursing and the helping professions: A user's guide*. Basingstoke: Palgrave Macmillan.
- Searle, J.R., 1969. *Speech acts: An essay in the philosophy of language*. Cambridge: Cambridge University Press.
- Tuckman, B.W., 1965. Developmental sequence in small groups. *Psychological Bulletin*, 63(6), pp. 384–399.
- Wooldridge, M.J., 2009. *An introduction to multiagent systems*. 2nd ed. Chichester: John Wiley & Sons.

Appendix A: e-portfolio artefacts and repositories

This appendix provides verifiable links to my team and individual contributions (discussion posts, reports, code repositories, and activity artefacts) to evidence collaboration.

| Units | Artefacts & Repositories |
|-----------|--|
| Main Page | https://m-kanuri.github.io/ |
| IA Module | https://m-kanuri.github.io/Module6.html |
| Unit 1 | 1) Main Page <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/07/29/IA-Unit1.html |

| | |
|--------|--|
| | <p>2) Collaborative Discussion 1 – Agent Based Systems</p> <ul style="list-style-type: none"> • Initial Post: https://m-kanuri.github.io/artefacts/IA-Unit01-Initial_Post.pdf <p>3) Module Wiki</p> <ul style="list-style-type: none"> • https://m-kanuri.github.io/artefacts/IA-Unit04-Module_Wiki.pdf |
| Unit 2 | <p>1) Main Page</p> <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/08/05/IA-Unit2.html <p>2) Collaborative Discussion 1 – Agent Based Systems</p> <ul style="list-style-type: none"> • Peer Response: https://m-kanuri.github.io/artefacts/IA-Unit01-Peer_Response.pdf |
| Unit 3 | <p>1) Main Page</p> <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/08/12/IA-Unit3.html <p>2) Collaborative Discussion 1 – Agent Based Systems</p> <ul style="list-style-type: none"> • Summary Post: https://m-kanuri.github.io/artefacts/IA-Unit01-SummaryPost.pdf |
| Unit 4 | <p>1) Main Page</p> <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/08/19/IA-Unit4.html <p>2) Module Wiki</p> <ul style="list-style-type: none"> • https://m-kanuri.github.io/artefacts/IA-Unit04-Module_Wiki.pdf (Continued till end of the module) |
| Unit 5 | <p>1) Main Page</p> <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/08/26/IA-Unit5.html <p>2) Collaborative Discussion 2 – Agent Communication Languages</p> <ul style="list-style-type: none"> • Initial Post: https://m-kanuri.github.io/artefacts/IA-Unit05-Initial_Post.pdf |
| Unit 6 | <p>1) Main Page</p> <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/09/02/IA-Unit6.html <p>2) Creating Agent Dialogues</p> <ul style="list-style-type: none"> • Activity Guidance: https://m-kanuri.github.io/artefacts/IA-Unit06-CreatingAgentDialogues.pdf <p>3) Collaborative Discussion 2 – Agent Communication Languages</p> <ul style="list-style-type: none"> • Peer Response: https://m-kanuri.github.io/artefacts/IA-Unit05-Peer_Response.pdf |

| | |
|---------|---|
| | 4) Development Team Project <ul style="list-style-type: none"> • https://m-kanuri.github.io/artefacts/IA-Unit06-DevelopmentTeamProjectReport.pdf |
| Unit 7 | 1) Main Page <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/09/09/IA-Unit7.html 2) Collaborative Discussion 2 – Agent Communication Languages <ul style="list-style-type: none"> • Summary Post: https://m-kanuri.github.io/artefacts/IA-Unit05-SummaryPost.pdf |
| Unit 8 | 1) Main Page <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/09/16/IA-Unit8.html 2) Parse Tree <ul style="list-style-type: none"> • https://m-kanuri.github.io/artefacts/IA-Unit08-Parse_Tree.pdf |
| Unit 9 | 1) Main Page <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/09/23/IA-Unit9.html 2) Collaborative Discussion 2 – Agent Communication Languages <ul style="list-style-type: none"> • Initial Post: https://m-kanuri.github.io/artefacts/IA-Unit09-Initial_Post.pdf |
| Unit 10 | 1) Main Page <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/09/30/IA-Unit10.html 2) Collaborative Discussion 2 – Agent Communication Languages <ul style="list-style-type: none"> • Peer Response: https://m-kanuri.github.io/artefacts/IA-Unit09-Peer_Response.pdf |
| Unit 11 | 1) Main Page: <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/10/07/IA-Unit11.html 2) Collaborative Discussion 2 – Agent Communication Languages <ul style="list-style-type: none"> • Summary Post: https://m-kanuri.github.io/artefacts/IA-Unit09-SummaryPost.pdf 3) Development Individual Project <ul style="list-style-type: none"> • Presentation: https://m-kanuri.github.io/artefacts/IA-Unit11-Development%20Individual%20Project-%20Presentation.pdf |

| | |
|---------|---|
| | <ul style="list-style-type: none"> • Transcript : https://m-kanuri.github.io/artefacts/IA-Unit11-Development%20Individual%20Project%20-%20Transcript%20Text.pdf • GitHub Repo: https://github.com/m-kanuri/ia_agent |
| Unit 12 | <p>1) Main Page:</p> <ul style="list-style-type: none"> • https://m-kanuri.github.io/module6/2025/10/14/IA-Unit12.html <p>2) Reflection: This document</p> <p>3) e-portfolio:</p> <ul style="list-style-type: none"> • https://m-kanuri.github.io/ (m-kanuri, 2024, updated 2025) • https://m-kanuri.github.io/Module6.html |

Appendix B: Unit 6 Team E Project Report

JIRA Board

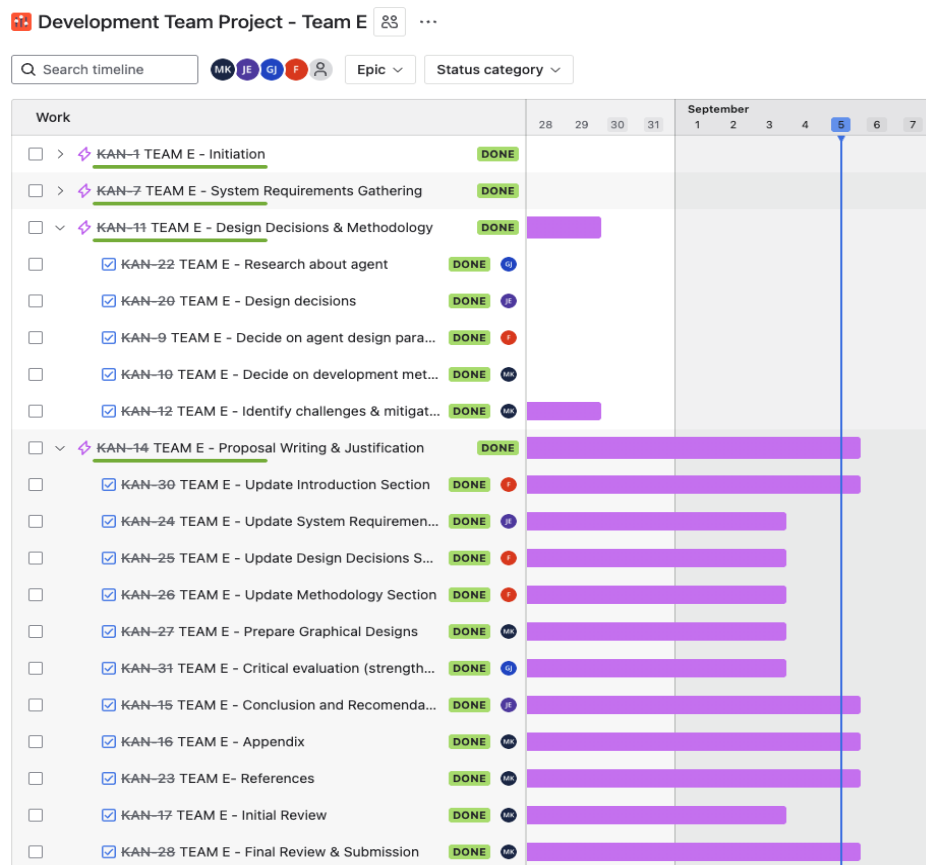


Figure B1: JIRA Board from Unit 6 Team E Project Report, adapted from Kanuri et al. (2025a) and Kanuri (2025b). Source: Author's own work.

GitHub Repository

The screenshot shows a GitHub repository named 'ia-group-e' (Public) by user 'Finlay-K'. The repository has 0 Watchers, 1 Fork, and 0 Stars. The main branch is selected, showing 1 Branch and 0 Tags. A search bar and 'Add file' button are visible. The repository contains a commit history table with files like 'project', 'proposal', '.gitignore', 'README.md', 'requirements.txt', and 'requirements_deps.txt'. The README is titled 'Digital Forensics Agent Part 1 - Group Project - Project Proposal' and describes the repository's purpose. A 'Project Structure' section shows a tree view of the files. The right sidebar includes 'About' (Intelligent Agents July 2025 A - Digital Forensics Project), 'Releases' (No releases published), 'Packages' (No packages published), 'Contributors' (3 contributors: m-kanuri, jaco-uoeo, Finlay-K), and 'Languages' (Python 100.0%).

| File | Commit Message | Time |
|-----------------------|---------------------------------|--------------|
| project | Initial commit | 2 months ago |
| proposal | Update benchmark_forensic_AI.py | last month |
| .gitignore | Initial commit | 2 months ago |
| README.md | Update README.md | 2 months ago |
| requirements.txt | Magic on Windows | 2 months ago |
| requirements_deps.txt | Magic on Windows | 2 months ago |

Digital Forensics Agent Part 1 - Group Project - Project Proposal

This repository contains benchmarking tools and prototype agents for the Digital Forensics Agent group project in the Intelligent Agents module. It is intended for experimentation with file-type detection, metadata extraction, and learning-based classification.

Project Structure

```
IA-GROUP-E/  
├── proposal/  
│   ├── benchmarks/  
│   │   └── mock_data/  
│   ├── agents/  
│   └── benchmark_runner.py
```

- # Shared experiments, benchmarks, and utilities
- # Performance tests for libraries and approaches
- # Mock data for shared testing
- # Agent prototypes
- # Benchmark runner helper with CLI support

Figure B2: GitHub repository contributions from Unit 6 Team E Project Report, adapted from Kanuri et al. (2025a) and Kanuri (2025b). Source: Author's own work.

Appendix C: Selected e-portfolio evidence

Collaborative Discussions

The screenshot shows a forum post titled "Collaborative Discussion 1: Agent Based Systems" on the University of Essex online platform. The post is a "Summary Post" by Linga Murthy Kanuri, dated Sunday, 17 August 2025, 5:43 PM. The post content includes an initial post summary and a list of peer responses. The interface includes a navigation bar with links like Home, Dashboard, My Modules, Student Services, Student Life, Student Voice, and Library. A search bar and a "Recent" dropdown are also visible. A "Knowledge Base" sidebar is on the right, and a "Chat to us!" button is at the bottom right.

University of Essex Online Home Dashboard My Modules Student Services Student Life Student Voice Library Recent Search LK

Collaborative Discussion 1: Agent Based Systems

Forum Reports

Summary Post

Summary Post Initial post Settings

Display replies in nested form

Summary Post

by Linga Murthy Kanuri - Sunday, 17 August 2025, 5:43 PM

In my initial post, I highlighted that adopting agent-based systems (ABS) is driven by three converging trends: ubiquity, interconnection, and intelligence (Wooldridge, 2009). These systems support distributed problem-solving, demonstrate autonomy, and evolve with organisational needs (Jennings and Wooldridge, 1998; Russell and Norvig, 2021). For Instance, ABS in a smart grid may automatically redirect energy to avoid blackouts, while trading bots make real-time choices about the market. I argued that ABS reduce human cognitive burden and enhance resilience by acting as 'people working together in software.'

The peer responses enriched this perspective.

- Jaafar emphasised emergent collective intelligence, noting that interacting agents often generate solutions beyond individual capacity (Bonabeau, 2002). He also linked ABS to organisational learning (Argote and Miron-Spektor, 2011).
- Building on this, Nasser recognised the benefits but cautioned that decentralisation raises challenges of coordination, trust, and accountability (Wooldridge, 2009; Mahela et al., 2020). He emphasised that autonomy introduces moral and legal challenges, particularly around liability and transparency (Florida and Cows, 2019).
- Martyna extended the discussion by focusing on human-agent collaboration. She contended that mixed-initiative teaming is beneficial in high-risk fields such as disaster management and healthcare (Klein et al., 2004). She underlined explainability as essential for trust and system integration (Gunning and Aha, 2019).

In my responses to peers, I engaged with these points.

- To Jaafar, I noted that autonomy can lead to conflicting objectives or unpredictable outcomes (Jennings and Wooldridge, 1998).
- I emphasised that ABS regulation will be greatly aided by governance frameworks like the EU AI Act (European Commission, 2024).
- To Martyna, I expanded on emergence using real-world analogies such as bird flocking and traffic flows (Bonabeau, 2002; Davidsson, Persson and Holmgren, 2007). I also questioned whether decentralisation, while improving resilience,

Chat to us!

The screenshot shows a forum post titled "Collaborative Discussion 2: Agent Communication Languages" on the University of Essex online platform. The post is a "Summary Post" by Linga Murthy Kanuri, dated Monday, 15 September 2025, 9:59 PM. The post content includes an initial post summary and a list of peer responses. The interface is similar to the first screenshot, with a navigation bar, search bar, and "Knowledge Base" sidebar.

University of Essex Online Home Dashboard My Modules Student Services Student Life Student Voice Library Recent Search LK

Collaborative Discussion 2: Agent Communication Languages

Forum Reports

Summary Post

Summary Post Summary Post Settings

Display replies in nested form

Summary Post

by Linga Murthy Kanuri - Monday, 15 September 2025, 9:59 PM

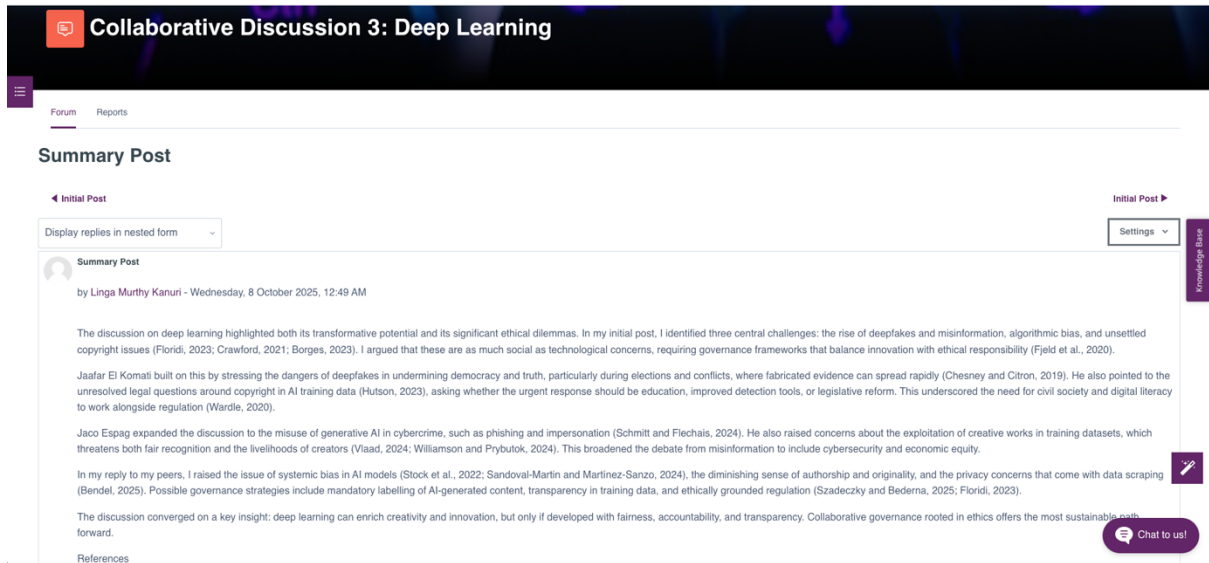
In my initial post, I discussed KQML (the Knowledge Query and Manipulation Language) as an agent communication language (ACL), comparing it to a method call in Python or Java. Semantic richness, interoperability and flexibility are the benefits of using ACLs because it is possible to communicate through semantically rich terms (Souza et al., 2016). However, they also present problems: wordiness, overhead of processing time and human efforts in creating and maintaining shared ontologies also become a great issue that should be solved properly (Köne et al., 2000). On the other hand, method calling is efficient and easy but suffers from lack of expressiveness and communication intention (Berna-Koes, Nourbakhsh & Sycara, 2004).

My colleagues contributed additional perspectives that expanded the conversation. Concerns raised in other work regarding ontology management are also evident, with domain specific ontologies and mediation affecting scalability and ACL's adding overhead (Jimenez-Ruiz et al., 2018). Jaafar stressed the significance of ontology alignment and called for more standardisation and automation support to deal with semantic mismatches (Mohammadi et al., 2017). He also suggested a middle way, where simple interactions could be supported by lightweight protocols and ACLs would take care of more complex exchanges (Suchanek, Abiteboul and Senellart, 2011). Ahmad emphasised semantic interoperability and advocated for middleware and modular ontologies as strategies to reduce overhead (Fatras, Ma and Jørgensen, 2022; Williams, 2025). He also related ACLs to knowledge engineering, emphasising that it "all boils down to having good feature representations necessary for efficient AI communication" (Souza et al., 2016).

In reaction, I explored these recommendations by examining hybrid protocols, ontology alignment techniques, and middleware optimisations. While such an architecture is useful in many regards, I concluded that ACLs and method invocation are not mutually exclusive possibilities; rather, they represent two points on a spectrum of potential designs, determined by system constraints, domain-specific requirements, and performance characteristics (Berna-Koes, Nourbakhsh and Sycara, 2004; Fatras, Ma and Jørgensen, 2022).

Reflecting on this conversation, the most vivid takeaway was how peers propelled the discussion beyond theoretical distinctions toward practical paths—such as ontology tools, hybrid methods, and middleware—that increase the real-world value

Chat to us!



Figures C1-C3: Screenshots showing the Summary post of the collaborative discussions

Module Wiki

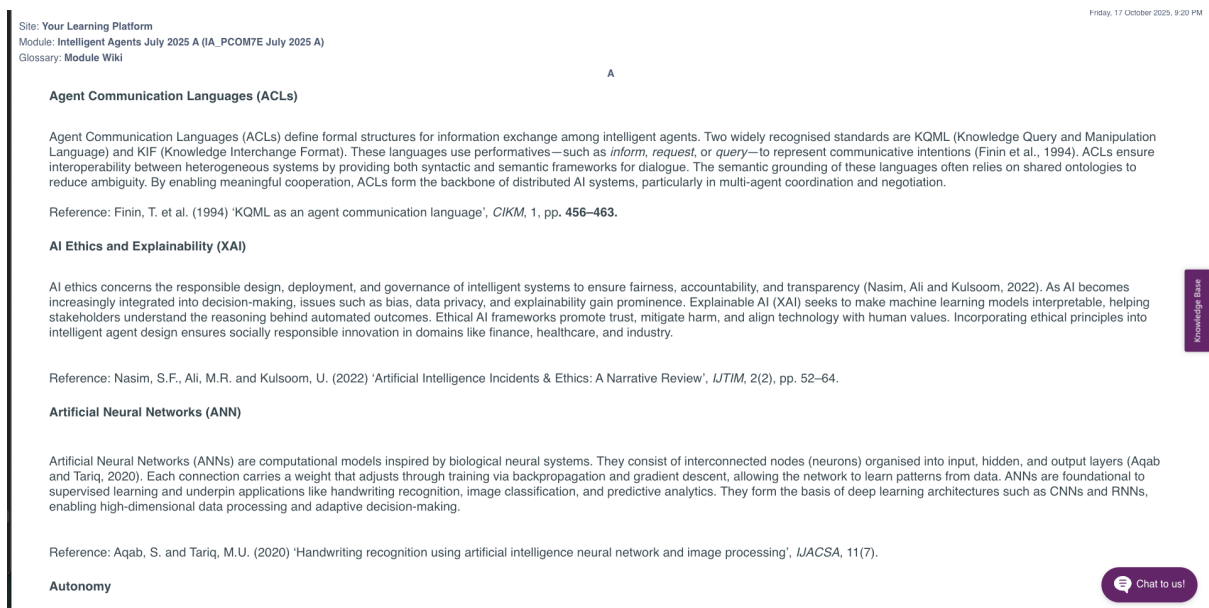


Figure C4: Screenshot showing the contribution to Module Wiki

Creating Parse Trees

The following figures illustrate my work from the *Parse Tree Creation* exercise, demonstrating syntax structure and parsing stages for different sample sentences.

1) The government raised interest rates.




Figure 1. Constituency parse tree for "The government raised interest rates." (Created using draw.io.)

2) The internet gives everyone a voice.




Figure 2. Constituency parse tree for "The internet gives everyone a voice." (Created using draw.io.)

3) The man saw the dog with the telescope.




Figure 3. Constituency parse tree for "The man saw the dog with the telescope." (Created using draw.io.)

Figures C5-7: Screenshots showing the Parse tree exercise answers

Deep Learning in Action

Overview of the Technology (What it Does)

This technology uses deep learning—typically convolutional neural networks (CNNs)—to analyse medical images such as X-rays, MRIs, CT scans and pathology slides (Esteva et al., 2017; Ardila et al., 2019). The aim is to detect, classify and quantify disease features that may be subtle or hard for the human eye to identify consistently. Applications include early cancer detection (e.g., lung nodules, skin lesions) and screening for diabetic retinopathy (Gulshan et al., 2016), with potential to improve reporting consistency and throughput, and—if evaluated across diverse populations—to support health equity (WHO, 2021).

Brief Synopsis of How it Works

Deep models require large, labelled datasets—typically thousands of expert-annotated images (Esteva et al., 2017).

- **Training:** The CNN learns spatial patterns linked to disease (e.g., texture/morphology separating malignant tumours from benign tissue).
- **Feature extraction:** Deep nets learn hierarchical features directly from data, avoiding manual feature engineering (LeCun, Bengio & Hinton, 2015).
- **Prediction:** In deployment, the model processes a new **unlabelled** scan and outputs a probability (e.g., likelihood of malignancy) to assist the clinician; many pipelines add **explanations** (e.g., Grad-CAM heatmaps) for transparency (Selvaraju et al., 2017).

Potential Socio-Technical Impacts (Ethics, Privacy, Fairness)

Using deep learning in high-stakes industries like health raises several important ethical and social issues.

Figure C8: Screenshot showing the Deep Learning in Action document and link shared in e-portfolio links.

GitHub Repository for Individual Project

The screenshot shows a GitHub repository named 'ia_agent' by user 'm-kanuri'. The repository is public and has 1 branch and 0 tags. The main branch is selected. The repository contains several files and folders, including 'TestEvidence', 'sample_data', 'src', 'tests', '.coveragerc', '.gitignore', 'README.md', 'requirements.txt', and 'transcript.md'. The README file is open, showing the title 'Intelligent Forensics Agent (Agent 1 Implementation)' and a description of the project. The README also includes a section for '1) Implementation Overview & Design Rationale'.

ia_agent Public

main 1 Branch 0 Tags

Go to file Add file Code

m-kanuri Add files via upload 303b2df · 5 days ago 88 Commits

| File/Folder | Commit Message | Time |
|------------------|---|------------|
| TestEvidence | Add files via upload | 5 days ago |
| sample_data | Initial: Agent 1 (safe discovery, content-first ID, SQLite, CL... | last week |
| src | Update database.py | last week |
| tests | Initial: Agent 1 (safe discovery, content-first ID, SQLite, CL... | last week |
| .coveragerc | Initial: Agent 1 (safe discovery, content-first ID, SQLite, CL... | last week |
| .gitignore | Clean repo: remove stray dirs; untrack venv/.coverage; ad... | last week |
| README.md | Update README.md | 5 days ago |
| requirements.txt | Initial: Agent 1 (safe discovery, content-first ID, SQLite, CL... | last week |
| transcript.md | Initial: Agent 1 (safe discovery, content-first ID, SQLite, CL... | last week |

README

Intelligent Forensics Agent (Agent 1 Implementation)

Practical implementation of my Unit 6 design for an **Intelligent Forensics Agent**.

Agent 1 (Implemented): safe discovery (read-only), **content-first** file identification, minimal metadata extraction, and storage in SQLite. **Scope Note:** Includes unit tests and a tiny, optional ML demo (logistic regression) for advisory "archive?" recommendations (Agent 2 groundwork).

python 3.11 tests unittest | coverage

1) Implementation Overview & Design Rationale

Figure C9: GitHub Repository from *Unit 11 Development Individual Project Presentation*, adapted from Kanuri (2025c). Source: Author's own work.

SQLite Screenshot

The screenshot shows the DB Browser for SQLite application interface. The title bar indicates the database path: `/Users/murthykanuri/Documents/ia_agent/agent.db`. The menu bar includes options like `New Database`, `Open Database`, `Write Changes`, `Revert Changes`, `Undo`, `Open Project`, `Save Project`, and `Attach Database`. The toolbar contains icons for file operations and execution. The main window is divided into three panes: `Database Structure`, `Browse Data`, and `Execute SQL`. The `Execute SQL` pane is active, showing a list of SQL queries with line numbers 1 through 16. The queries include comments and SQL statements for counting rows, MIME breakdown, largest files, and duplicate content by hash. Below the queries, the results of the fifth query are displayed in a table with two columns: `mime` and `files`. The results show five rows of data. At the bottom, a status message indicates that the execution finished without errors and that 5 rows were returned in 4ms.

```
1  -- how many rows?
2  SELECT COUNT(*) AS rows FROM files;
3
4  -- MIME breakdown
5  SELECT mime, COUNT(*) AS files
6  FROM files GROUP BY mime ORDER BY files DESC;
7
8  -- largest files
9  SELECT path, size
10 FROM files ORDER BY size DESC LIMIT 10;
11
12 -- duplicate content by hash
13 SELECT sha256, COUNT(*) AS n
14 FROM files GROUP BY sha256 HAVING n > 1 ORDER BY n DESC;
15
16 Select * from files;
```

| | mime | files |
|---|--------------------------|-------|
| 1 | text/plain | 2 |
| 2 | application/octet-stream | 1 |
| 3 | application/pdf | 1 |
| 4 | image/jpeg | 1 |
| 5 | text/csv | 1 |

Execution finished without errors.
Result: 5 rows returned in 4ms
At line 5:
SELECT mime, COUNT(*) AS files
FROM files GROUP BY mime ORDER BY files DESC;

Figure C10: SQLite Screenshot from *Unit 11 Development Individual Project Presentation*, adapted from Kanuri (2025c). Source: Author's own work.

Appendix D: Personal Skills Matrix and Action Plan (PDP)

| SKILL | LEVEL | EVIDENCE | ACTION PLAN |
|--------------------------------|--------------|--|--|
| TIME MANAGEMENT | Expert | Attending all the seminar sessions, meeting all deadlines. Over 25 years of professional experience working on IT projects. | Continue to manage the time effectively, delivering work on schedule and maintaining quality results. |
| CRITICAL THINKING AND ANALYSIS | Proficient | Evaluated the advantages and limitations of hybrid agent architectures. Analysed case studies on AI ethics and decision-making in autonomous systems. | Strengthen by applying structured evaluation frameworks (e.g., SWOT or ethical impact analysis) to future AI design tasks. |
| COMMUNICATION AND LITERACY | Proficient | Contributed actively to all the collaborative discussions. Moderated the Team E group project as a project Manager | Improve by producing concise technical summaries and enhancing clarity in collaborative documentation. |
| IT AND DIGITAL | Advanced | Developed, tested, and maintained code for intelligent agent systems using Python and GitHub. | Continue exploring advanced AI frameworks (e.g., TensorFlow, spaCy) and apply version control best practices in future development. |
| NUMERACY | Intermediate | Applied quantitative reasoning to evaluate the accuracy and loss functions of ANN models in the <i>Deep Learning in Action</i> activity. Used data outputs and confusion matrices to assess model performance. | Enhance numerical analysis by studying statistical validation methods such as cross-validation and standard deviation in model testing. Incorporate data visualisation to communicate performance metrics clearly. |
| RESEARCH | Advanced | Conducted literature review on AI ethics and evolution; correctly applied Harvard referencing in e-portfolio reflections. | Continue systematic literature searches; summarise two new AI-ethics articles each month for ongoing professional development. |
| INTERPERSONAL | Proficient | Collaborated effectively within the Unit 6 development team and peer-review tasks; demonstrated accountability and adaptability. | Take leadership or coordination roles in future projects; mentor peers to |

| | | | |
|-------------------|------------|--|---|
| | | | strengthen facilitation and negotiation skills. |
| PROBLEM SOLVING | Advanced | Diagnosed coding issues during the team project and individual development project. Integrated feedback from peers to improve functionality. | Formalise a debugging checklist (repro steps, logs, tests, fix, verify); increase test coverage (unit/integration) for agent modules; run post-mortems after defects to capture lessons learned in the e-portfolio. |
| ETHICAL AWARENESS | Proficient | Reflected on ethical implications of intelligent agents in discussions and portfolio; considered bias and accountability. | Expand understanding through IEEE and ACM ethical AI frameworks; integrate ethics checklists in future system designs. |

Appendix E: Skills Matrix

- Consolidated skills matrix can be found here: <https://github.com/m-kanuri/m-kanuri.github.io/blob/main/artefacts/Skills%20Matrix.xlsx>